

Redis Enumeration - Complete Theory & Practice Notes

Introduction to Databases & Redis

Database Fundamentals

Databases are collections of organized information that can be easily accessed, managed and updated. In most environments, database systems are critical because they communicate information related to:

- Sales transactions
- Product inventory
- Customer profiles
- Marketing activities

Types of Databases

There are different types of databases, and **Redis** is an '**in-memory**' database.

In-memory databases:

- Rely essentially on primary memory (RAM) for data storage
- Database is managed in the RAM of the system
- Contrast to databases that store data on disk or SSDs
- Primary memory is significantly faster than secondary memory
- Data retrieval time is very small, offering efficient & minimal response times

Redis Usage Scenarios

Common Implementation Pattern:

1. Website needs to display prices on front page
2. Website first checks if prices are in Redis cache
3. If not found, checks traditional database (MySQL/MongoDB)
4. When value is loaded from database, it's stored in Redis for shorter period (seconds/minutes/hours)
5. Handles similar requests during that timeframe from cache

Benefits for High-Traffic Sites:

- Much faster retrieval for majority of requests
- Stable long-term storage in main database
- Reduced load on primary database

What is Redis?

Redis (REmote DIctionary Server) is an open-source advanced NoSQL key-value data store used as:

- **Database:** Primary data storage
- **Cache:** Temporary fast storage
- **Message broker:** Communication between services

Core Architecture

Server Component

- Redis runs as server-side software
- Core functionality is in server component
- Server listens for connections from clients
- Accepts connections programmatically or through CLI

Database Storage

- Data stored in dictionary format with key-value pairs
- Database stored in server's RAM for fast data access
- Redis writes database contents to disk at varying intervals
- Provides persistence as backup in case of failure
- Typically used for short-term storage of data needing fast retrieval

The CLI (Command-Line Interface)

- Powerful tool giving complete access to Redis data and functionalities
- Essential for developing software/tools that interact with Redis
- Provides direct database manipulation capabilities

Enumeration Methodology

Phase 1: Target Verification

Connectivity Check: Verify target availability and connection quality

```
bash
```

```
ping <target_IP>
```

```
# Run until satisfied with connection (2-3 replies sufficient)
```

```
# Sometimes getting snippet/overview is more time-efficient than detailed reports
```

Phase 2: Network Reconnaissance

Port Scanning: Identify open services and versions

```
bash

nmap -p- -sV <target_IP>
# -sV: Probe open ports to determine service/version info
```

Expected Redis Indicators:

- Port 6379/tcp open
- Service: Redis key-value store
- Version information displayed

Phase 3: Redis-Specific Enumeration

Tool Installation

```
bash

# Install redis-cli utility
apt install redis-tools
# Alternative: netcat can connect but redis-cli is more convenient
```

Connection Establishment

```
bash

redis-cli -h <target_IP>
# Should display prompt: target_IP:6379>
```

Redis CLI Reference

Essential Connection Options

- `-h <hostname>`: Server hostname (default: 127.0.0.1)
- `-p <port>`: Server port (default: 6379)
- `-t <timeout>`: Server connection timeout in seconds
- `-a <password>`: Password for authentication
- `--user <username>`: ACL style authentication with username
- `-u <uri>`: Server URI format: redis://user:password@host:port/dbnum
- `-n <db>`: Database number to select
- `-c`: Enable cluster mode

- `--tls`: Establish secure TLS connection

Advanced Options

- `-r <repeat>`: Execute specified command N times
- `-i <interval>`: Wait interval between repeated commands
- `--raw`: Use raw formatting for replies
- `--csv`: Output in CSV format
- `--json`: Output in JSON format
- `--scan`: List all keys using SCAN command
- `--pattern <pat>`: Key pattern for scanning operations
- `--bigkeys`: Sample keys looking for large elements
- `--hotkeys`: Sample keys looking for frequently accessed keys

Core Redis Commands & Theory

Information Gathering

```
bash
```

```
info
```

Purpose: Returns comprehensive information and statistics about Redis server **Output Sections:**

- **Server:** Version, OS, configuration files, process info
- **Clients:** Connected clients, buffer information
- **Memory:** Usage statistics, fragmentation, allocator info
- **Persistence:** RDB and AOF status, backup information
- **Stats:** Connection counts, command processing, network I/O
- **Replication:** Master/slave configuration
- **CPU:** System and user CPU usage
- **Keyspace:** Database statistics including key counts

Database Selection & Navigation

```
bash
```

```
select <database_number>
```

Theory: Redis supports multiple logical databases (default 16, numbered 0-15) **Usage:** Switch between databases for organization/isolation

Key Discovery

```
bash
keys *
keys <pattern>
```

Theory:

- Lists keys matching pattern (* = all keys)
- Can be slow on large databases (use SCAN for production)
- Pattern matching supports glob-style patterns

Data Retrieval

```
bash
get <key_name>
type <key_name>
dbsize
```

Functions:

- `get`: Retrieve string value associated with key
- `type`: Determine data type (string, list, set, hash, etc.)
- `dbsize`: Count total keys in current database

Practical Enumeration Example

Target Analysis

Network Scan Results:

```
$ nmap -p- -sV 10.129.136.187
PORT      STATE SERVICE VERSION
6379/tcp  open  redis    Redis key-value store 5.0.7
```

Interpretation:

- Only port 6379 open (typical Redis configuration)
- Redis version 5.0.7 identified
- No additional services running

Connection & Initial Reconnaissance

```
bash  
  
redis-cli -h 10.129.136.187  
10.129.136.187:6379>
```

Server Information Analysis

```
bash  
  
10.129.136.187:6379> info  
# Server  
redis_version:5.0.7  
os:Linux 5.4.0-77-generic x86_64  
tcp_port:6379  
config_file:/etc/redis/redis.conf  
  
# Memory  
used_memory_human:839.48K  
total_system_memory_human:1.94G  
  
# Keyspace  
db0:keys=4,expires=0,avg_ttl=0
```

Key Findings:

- Redis 5.0.7 on Linux system
- Low memory usage (839KB) indicates small dataset
- Database 0 contains 4 keys with no expiration set
- No authentication required (security issue)

Database Exploration

```
bash  
  
10.129.136.187:6379> select 0 # Select database 0  
10.129.136.187:6379> keys * # List all keys  
1) "flag"  
2) "stor"  
3) "numb"  
4) "temp"  
  
10.129.136.187:6379> get flag # Retrieve flag value  
"03e1d2b376c37ab3f5319922053953eb"
```

Security Implications & Vulnerabilities

Common Redis Security Issues

1. **No Authentication:** Default installations often lack password protection
2. **Network Exposure:** Redis accessible from internet without firewall
3. **Privilege Escalation:** Redis running with elevated privileges
4. **Data Exposure:** Sensitive information stored in plaintext
5. **Command Injection:** Unsafe handling of user input in applications

Enumeration Indicators

- Direct connection without credentials
- `info` command execution successful
- Access to key listing and data retrieval
- Server configuration details exposed

Real-World Attack Scenarios

- **Data Exfiltration:** Extracting cached user sessions, API keys
- **Configuration Theft:** Application settings, database credentials
- **Cache Poisoning:** Injecting malicious data into cache
- **Lateral Movement:** Using Redis as pivot point in network

Enumeration Best Practices

Systematic Approach

1. **Start with `info`:** Comprehensive server overview
2. **Analyze keyspace:** Identify available databases and key counts
3. **Database enumeration:** Check each database systematically
4. **Pattern-based searching:** Use specific patterns for targeted discovery
5. **Data type analysis:** Understand data structures before extraction

Performance Considerations

- **Large databases:** Use `SCAN` instead of `KEYS *`
- **Pattern matching:** Specific patterns reduce response time
- **Batch operations:** Group related commands for efficiency

Documentation

- Record Redis version and configuration
- Document database structure and key naming conventions
- Note security configurations and access controls
- Catalog sensitive data discovered

Advanced Enumeration Techniques

Database Mapping

```
bash

info keyspace           # Quick database overview
scan 0 match *pattern* # Efficient key discovery
```

Data Structure Analysis

```
bash

type <key>           # Identify data type
strlen <key>          # String Length
llen <key>            # List Length
hlen <key>            # Hash field count
```

Configuration Analysis

```
bash

config get "*"        # Retrieve configuration parameters
client list            # Show connected clients
```

This comprehensive approach ensures thorough Redis enumeration while understanding the underlying technology and security implications.